



education

phd | cs

uc berkeley (bair) | '17-'22

research: interpretable ai

advisor: bin yu

bs | cs & math

university of virginia | '14-'17

double major

skills

llms • interpretability

pre-training • post-training

data science • causal inference

healthcare • neuroscience

teaching

berkeley | fall 2019

artificial intelligence: cs 188

berkeley | summer 2018

machine learning: cs 189/289

lectures to class of 80+ students

service

mentorship

cbai safety mentoring '25-'26

bair undergrad mentoring '18-'22

volunteering

basis education volunteering '19-'22

computer literacy volunteering '15-'17

organizer / area chair

iclr trustml workshop '26

ml4h '25

xxai workshop, ml4h '24

reviewer

iclr '27 • iclr, icml, neurips '26

arr, iclr, icml, iccv, neurips, pnas '25

iclr, icml, neurips '24

neurips '23 • acl '22

iclr, cvpr, aaai, neurips '21 • neurips '20

awards

berkeley grad slam semifinalist '19, '22

pdsoros fellowship finalist '19

outstanding teaching award '18

uva rader research award '17

uva undergrad symposium winner '17

icpc regional qualification '14-'16

1st place microsoft code jam '16

rodman scholarship '14-'17

experience

voyage ai (acquired by mongodb)

staff research scientist | summer '26 - present

- pre-training and adapting large embedding models

microsoft research

senior research scientist (deep learning group) | summer '22 - now

- researched interpretability & efficiency of llms
- built foundation models for neuroscience & medicine
- managed small group of phd students & interns

health tech

paige ai | research scientist | summer '21 - summer '22

- interpretable foundation models in digital pathology

response4life | volunteer data scientist | spring '20

- helped develop, integrate, and deploy models to forecast covid-19 severity

pacmed ai | healthcare ml intern | summer '19

- developed interpretable, tabular ai models for healthcare

phd

berkeley | interpretable ai research (bin yu group) | fall '17 - spring '22

- developed post-hoc interpretation methods for neural networks
- developed interpretable models in medicine, biology, and computer vision

aws | ai fairness intern (pietro perona group) | summer '20

- researched algorithmic bias with causal matching using GANs

meta ai | computer vision intern | summer '17

- investigated unsupervised deep learning for segmentation of satellite imagery

undergrad

hhmi | computer vision intern (srini turaga group) | summer '14, '15, '16

uva | graphical models research (yanjun qi group) | fall '16 – spring '17

uva | comp. neuroscience research (william levy group) | fall '14 - fall '16

selected works

neuroscience

- augmenting interpretable models with llms: **cs** et al. *nature comms*, '23
- generative causal testing with llms: antonello*, **cs***, et al. *nature neuro*, '26
- interpretable embeddings by asking llms questions benara*, **cs***, et al. *neurips*, '24

healthcare

- human-ai co-design for clinical models: feng*, kothari*, ..., **cs**. *npj digital med* '26
- interpretable models via autoresearch: **cs** et al. *neurips* '26 ; 1.6k+ Github stars
- fast interpretable greedy-tree sums: tan*, **cs***, nasseri*, agarwal* et al. *pnas* '22

general interpretability

- post-hoc attention steering for llms: zhang, **cs**, et al. *iclr* '24
- adaptive wavelet distillation from dnns: ha, **cs**, et al. *neurips* '21
- hierarchical interpretations for dnn predictions: **cs***, murdoch*, & yu. *iclr* '19